Robust Fault-Tolerant Policies in Reinforcement Learning

A Formal Theory of Policy Redundancy, Adversarial Fault Tolerance, and the Hand-in-Hand Nature of Adaptivity and Robustness

Mohsen Dirbaz, PhD

Abstract

We develop a formal theory of fault-tolerant policies in reinforcement learning. Let $\mathcal{M} = (\mathcal{S}, \mathcal{A}, P, R, \gamma)$ be a Markov decision process, $\pi : \mathcal{S} \to \mathcal{A}$ a policy, and $\mathcal{F}$ a fault model describing the perturbations that can occur to the policy's execution. We define a robust fault-tolerant policy as one that preserves a constitutional performance guarantee under every fault in $\mathcal{F}$. We prove a policy redundancy theorem: robust fault tolerance against a fault class $\mathcal{F}$ requires policy redundancy at least equal to the action-space entropy of $\mathcal{F}$, and the bound is tight for independent faults. We prove a recoverability characterization: a policy is recoverable from a fault if and only if the fault's affected state-action cone is closed and the recovery policy's domain contains the faulty successor. We prove a hand-in-hand theorem: fault tolerance and adaptivity are structurally coupled—the same policy redundancy that provides fault tolerance also provides the capacity for exploration and adaptation, and vice versa. The coupling is exact: the exploration entropy of a policy is bounded below by its fault-tolerance redundancy and above by its total stochasticity. We distinguish four fault models—observation, action, transition, and reward—by the component of the MDP they perturb. The framework connects to the synthesis through the learning and stochasticity structures but addresses a question the first synthesis does not: how a policy remains safe and adaptive under adversarial perturbation.

Program status. The mathematical content of this article is fully developed and self-contained. Its empirical content—the claim that it applies to specific reinforcement learning systems—remains a matter for independent validation. Field significance and novelty are empirical bibliographic questions that must survive adversarial prior-art audit.

Contents

I Foundations **3**

1 Introduction **3**
1.1 The surprising coupling 3
1.2 Relation to the program 4
1.3 Contributions 4
1.4 Falsification strategy 4

2 Preliminaries **5**
2.1 Markov decision processes 5
2.2 Constitutional constraints in RL 5
2.3 Policy entropy 6

II Fault Models for RL **6**

3 The Four Fault Models **6**
3.1 Observation faults . . . 6
3.2 Action faults . . . 7
3.3 Transition faults . . . 7
3.4 Reward faults . . . 7
3.5 The joint fault model . . . 8

III Robust Policies **8**

4 Policy Redundancy **8**
4.1 Definition . . . 8
4.2 Policy redundancy types . . . 9

5 Constitutional Robustness **9**
5.1 Definition . . . 9
5.2 The robust-value frontier . . . 10

IV Policy Redundancy Theorem **10**

6 Statement **10**
6.1 Action-space entropy of the fault class . . . 10
6.2 The theorem . . . 10

7 Tightness **11**
7.1 The bound is tight for independent action faults . . . 11

V Recoverability Characterization **11**

8 The Affected State-Action Cone **12**
8.1 Definition . . . 12
8.2 Recoverability . . . 12

9 Discussion **13**
9.1 The role of closure . . . 13

VI The Hand-in-Hand Theorem **13**

10 The Coupling **13**
10.1 Exploration entropy . . . 13
10.2 The hand-in-hand theorem . . . 14
10.3 The sandwich interpretation . . . 14

11 Consequences **14**
11.1 Pareto frontier of safety-adaptivity tradeoffs 14
11.2 Implications for policy design 15
11.3 The surprising implication 15

VII Composition and Hierarchy **15**

12 Composition of Robust Policies **15**
12.1 Composition 15
12.2 Closure of the hand-in-hand property 16

13 Hierarchical Robust Policies **16**
13.1 Hierarchy 16
13.2 Inter-level robustness 16

VIII Discussion **17**

14 The Hand-in-Hand Extension of the Synthesis **17**
14.1 The robustness-adaptivity determinant 17
14.2 Policy redundancy theorem 17
14.3 Hand-in-hand theorem 17
14.4 Four fault models 17

15 Relation to Prior Work **17**
15.1 Robust reinforcement learning 17
15.2 Exploration in reinforcement learning 18
15.3 Safe reinforcement learning 18

16 Limitations **18**

17 Open Problems **18**

18 Conclusion **19**

Part I
Foundations

1 Introduction

1.1 The surprising coupling

Fault tolerance and adaptivity are usually treated as separate engineering goals. Fault tolerance is about safety: keeping the system within its admissible set under perturbation. Adaptivity is about learning: exploring the environment to discover better policies. The two are often in tension: a conservative policy is safe but not adaptive; an exploratory policy is adaptive but not safe.

This article argues that the tension is superficial. The two are structurally coupled: the same policy property that provides fault tolerance also provides the capacity for exploration and adaptation. We call this the *hand-in-hand* property, and we prove it as a theorem.

The intuition is simple. A policy that can tolerate a fault in one action must have alternative actions available. The alternatives that provide fault tolerance also provide the exploration capacity. The exploration capacity that provides adaptivity also provides the fault-tolerance margin. The two are the same structural resource viewed from different angles.

The formal result is that the exploration entropy of a policy is bounded below by its fault-tolerance redundancy and above by its total stochasticity. This sandwich theorem makes the coupling exact.

1.2 Relation to the program

This article extends the second bundle in a direction the previous articles do not cover. The fault-tolerance article (Article 31, general version) developed the general theory of fault tolerance in adaptive systems. The present article specializes that theory to reinforcement learning policies and proves the hand-in-hand theorem.

The connection to the synthesis (Article 20) is through the learning determinant (Article 21) and the stochasticity determinant (Article 30). The learning determinant specifies how the system acquires structure; the stochasticity determinant specifies how the system behaves under noise. The present article identifies the structural feature that couples the two: policy redundancy.

1.3 Contributions

1. We define the fault-tolerant policy: a policy that preserves a constitutional performance guarantee under faults.
2. We define policy redundancy: the excess action-space capacity beyond what is needed for the nominal policy.
3. We prove the policy redundancy theorem: robust fault tolerance requires redundancy at least equal to the action-space entropy of the fault class.
4. We prove the tightness of the bound for independent faults.
5. We define the recoverability of a policy and prove the recoverability characterization.
6. We prove the hand-in-hand theorem: exploration entropy is bounded below by fault-tolerance redundancy and above by total stochasticity.
7. We distinguish four fault models: observation, action, transition, and reward.
8. We establish the composition of robust policies and the closure of the hand-in-hand property under composition.
9. We characterize the Pareto frontier of safety-adaptivity tradeoffs.

1.4 Falsification strategy

The policy redundancy theorem fails if there exists a robust fault-tolerant policy with redundancy below the action-space entropy of the fault class. The hand-in-hand theorem fails if there exists a

policy with exploration entropy below its fault-tolerance redundancy or above its total stochasticity. The recoverability characterization fails if there exists a recoverable fault whose affected cone is not closed.

2 Preliminaries

2.1 Markov decision processes

Definition 2.1 (Markov decision process)**.** *A* Markov decision process *(MDP) is a tuple*

$$\mathcal{M} = (\mathcal{S}, \mathcal{A}, P, R, \gamma)$$

where:

- $\mathcal{S}$ *is the state space;*
- $\mathcal{A}$ *is the action space;*
- $P : \mathcal{S} \times \mathcal{A} \to \mathcal{P}(\mathcal{S})$ *is the transition kernel;*
- $R : \mathcal{S} \times \mathcal{A} \to \mathbb{R}$ *is the reward function;*
- $\gamma \in [0, 1)$ *is the discount factor.*

Definition 2.2 (Policy)**.** *A* deterministic policy *is a function* $\pi : \mathcal{S} \to \mathcal{A}$*. A* stochastic policy *is a function* $\pi : \mathcal{S} \to \mathcal{P}(\mathcal{A})$ *assigning a distribution over actions to each state.*

Definition 2.3 (Value function)**.** *The* value function *of a policy* π *is*

$$V^\pi(s) = \mathbb{E}_\pi \left[\sum_{t=0}^{\infty} \gamma^t R(s_t, a_t) \mid s_0 = s \right].$$

Definition 2.4 (Optimal policy)**.** *The* optimal policy *is* $\pi^* = \arg\max_\pi V^\pi(s)$ *for all* $s \in \mathcal{S}$*.*

Proposition 2.5. *The value function satisfies the Bellman equation:*

$$V^\pi(s) = R(s, \pi(s)) + \gamma \sum_{s'} P(s' \mid s, \pi(s)) V^\pi(s').$$

Proof. By the law of total expectation. □

2.2 Constitutional constraints in RL

Definition 2.6 (Constitutional admissible set)**.** *The* constitutional admissible set *is a subset* $\mathcal{S}_{\text{adm}} \subseteq \mathcal{S}$ *of states satisfying the system's constitutional constraints.*

Definition 2.7 (Constitutional policy)**.** *A policy* π *is* constitutional *if it never transitions from* $\mathcal{S}_{\text{adm}}$ *to* $\mathcal{S} \setminus \mathcal{S}_{\text{adm}}$*:*

$$s \in \mathcal{S}_{\text{adm}} \implies P(s' \mid s, \pi(s)) = 0 \text{ for all } s' \notin \mathcal{S}_{\text{adm}}.$$

Definition 2.8 (Constitutional value)**.** *The* constitutional value *of a constitutional policy* π *is*

$$V_{\mathcal{C}}^\pi(s) = \mathbb{E}_\pi \left[\sum_{t=0}^{\infty} \gamma^t R(s_t, a_t) \mid s_0 = s, s_t \in \mathcal{S}_{\text{adm}} \text{ for all } t \right].$$

Proposition 2.9. *The constitutional value is well-defined for constitutional policies.*

Proof. The policy never leaves $\mathcal{S}_{\text{adm}}$, so the conditional expectation is over trajectories entirely within $\mathcal{S}_{\text{adm}}$. □

Remark 2.10. *Constitutional policies are the RL analogue of the constitutional transitions of Article 1. The admissible set replaces the feasible set; the constitutional value replaces the nominal value.*

2.3 Policy entropy

Definition 2.11 (Policy entropy)**.** *The* policy entropy *of a stochastic policy π at state s is*

$$H(\pi(\cdot \mid s)) = -\sum_{a \in \mathcal{A}} \pi(a \mid s) \log \pi(a \mid s).$$

Definition 2.12 (Total policy entropy)**.** *The* total policy entropy *is*

$$H(\pi) = \mathbb{E}_{s \sim \mu}[H(\pi(\cdot \mid s))]$$

where μ is the state visitation distribution.

Proposition 2.13. *The total policy entropy is non-negative.*

Proof. The entropy of any distribution is non-negative. □

Remark 2.14. *The policy entropy measures the stochasticity of the policy. A deterministic policy has zero entropy; a uniform policy has maximal entropy.*

Part II
Fault Models for RL

3 The Four Fault Models

3.1 Observation faults

Definition 3.1 (Observation fault)**.** *An* observation fault *is a perturbation of the state observation:*

$$\phi_{\text{obs}} : \mathcal{S} \to \mathcal{S}, \quad s \mapsto \phi_{\text{obs}}(s).$$

An observation fault corrupts the state the policy sees. The policy acts on $\phi_{\text{obs}}(s)$ instead of s.

Definition 3.2 (Observation-robust policy)**.** *A policy π is* observation-robust *against a fault class $\mathcal{F}_{\text{obs}}$ if*

$$V_{\mathcal{C}}^{\pi \circ \phi}(s) \geq V_{\mathcal{C}}^{\pi}(s) - \epsilon$$

for all $\phi \in \mathcal{F}_{\text{obs}}$, where $\pi \circ \phi(s) = \pi(\phi(s))$ and $\epsilon > 0$ is the tolerance.

Proposition 3.3. *Observation robustness requires the policy to be insensitive to the fault's perturbations.*

Proof. If the policy changes its action in response to the fault, the value may decrease by more than ϵ. □

Remark 3.4. *Observation faults model sensor noise, adversarial perturbations, or distribution shift. They are the standard fault model in robust RL.*

3.2 Action faults

Definition 3.5 (Action fault)**.** *An* action fault *is a perturbation of the executed action:*

$$\phi_{\text{act}} : \mathcal{A} \to \mathcal{A}, \quad a \mapsto \phi_{\text{act}}(a).$$

An action fault corrupts the action the policy executes. The executed action is $\phi_{\text{act}}(\pi(s))$ instead of $\pi(s)$.

Definition 3.6 (Action-robust policy)**.** *A policy π is* action-robust *against a fault class $\mathcal{F}_{\text{act}}$ if*

$$V_{\mathcal{C}}^{\phi \circ \pi}(s) \geq V_{\mathcal{C}}^{\pi}(s) - \epsilon$$

for all $\phi \in \mathcal{F}_{\text{act}}$.

Proposition 3.7. *Action robustness requires the policy to be robust to action-space perturbation.*

Proof. If the policy's action is perturbed, the resulting value may deviate. The policy must be designed so that the deviation is bounded. □

Remark 3.8. *Action faults model actuator failures, control errors, or adversarial action perturbation. They are the standard fault model in fault-tolerant control.*

3.3 Transition faults

Definition 3.9 (Transition fault)**.** *A* transition fault *is a perturbation of the transition kernel:*

$$\phi_{\text{trans}} : P \to P', \quad P'(s' \mid s, a) = P(s' \mid s, a) + \delta(s, a, s').$$

A transition fault corrupts the environment's dynamics. The actual transition differs from the nominal transition.

Definition 3.10 (Transition-robust policy)**.** *A policy π is* transition-robust *against a fault class $\mathcal{F}_{\text{trans}}$ if*

$$V_{\mathcal{C}}^{\pi,P'}(s) \geq V_{\mathcal{C}}^{\pi,P}(s) - \epsilon$$

for all $P' \in \mathcal{F}_{\text{trans}}$.

Proposition 3.11. *Transition robustness requires the policy to be robust to dynamics perturbation.*

Proof. If the environment's dynamics change, the value may decrease. The policy must be designed so that the decrease is bounded. □

Remark 3.12. *Transition faults model environmental change, distribution shift, or adversarial dynamics perturbation. They are the standard fault model in robust control and domain adaptation.*

3.4 Reward faults

Definition 3.13 (Reward fault)**.** *A* reward fault *is a perturbation of the reward function:*

$$\phi_{\text{rew}} : R \to R', \quad R'(s, a) = R(s, a) + \delta(s, a).$$

A reward fault corrupts the reward signal. The policy's observed rewards differ from the nominal rewards.

Definition 3.14 (Reward-robust policy)**.** *A policy π is* reward-robust *against a fault class $\mathcal{F}_{\text{rew}}$ if*

$$V_{\mathcal{C}}^{\pi,R'}(s) \geq V_{\mathcal{C}}^{\pi,R}(s) - \epsilon$$

for all $R' \in \mathcal{F}_{\text{rew}}$.

Proposition 3.15. *Reward robustness requires the policy to be robust to reward perturbation.*

Proof. If the reward signal is corrupted, the policy's optimization target changes. The policy must be designed so that the value decrease is bounded. □

Remark 3.16. *Reward faults model reward hacking, misspecified objectives, or adversarial reward manipulation. They are the standard fault model in inverse RL and reward-robust RL.*

3.5 The joint fault model

Definition 3.17 (Joint fault model)**.** *The* joint fault model *is*

$$\mathcal{F}_{\text{joint}} = \mathcal{F}_{\text{obs}} \times \mathcal{F}_{\text{act}} \times \mathcal{F}_{\text{trans}} \times \mathcal{F}_{\text{rew}}.$$

Proposition 3.18. *Joint robustness is stronger than robustness against any individual fault model.*

Proof. A policy robust against the joint model is robust against each individual model. The converse is not true in general. □

Remark 3.19. *Joint robustness is the strongest fault-tolerance requirement. It is the standard for safety-critical RL systems.*

Part III
Robust Policies

4 Policy Redundancy

4.1 Definition

Definition 4.1 (Policy redundancy)**.** *The* policy redundancy *of a constitutional policy π at state s is*

$$r(\pi, s) = \log |\mathcal{A}_{\text{adm}}(s)| - H(\pi(\cdot \mid s))$$

where $\mathcal{A}_{\text{adm}}(s)$ is the set of admissible actions at state s and $H(\pi(\cdot \mid s))$ is the policy's entropy at s.

The policy redundancy measures the excess action-space capacity beyond what the policy's nominal entropy uses. It is the structural reserve the policy has to absorb faults.

Proposition 4.2. *The policy redundancy is non-negative.*

Proof. $H(\pi(\cdot \mid s)) \leq \log |\mathcal{A}_{\text{adm}}(s)|$ by the maximum-entropy bound. □

Proposition 4.3. *The policy redundancy is maximized when the policy is deterministic.*

Proof. If the policy is deterministic, $H(\pi(\cdot \mid s)) = 0$, so $r(\pi, s) = \log |\mathcal{A}_{\text{adm}}(s)|$. □

Remark 4.4. *The policy redundancy is the RL analogue of the system redundancy of Article 31. It is the structural reserve the policy has to tolerate faults.*

4.2 Policy redundancy types

Definition 4.5 (Action redundancy)**.** Action redundancy *is the number of alternative actions available at each state.*

Definition 4.6 (Value redundancy)**.** Value redundancy *is the gap between the value of the optimal action and the value of the second-best action.*

Definition 4.7 (State redundancy)**.** State redundancy *is the number of alternative states from which the same action is admissible.*

Proposition 4.8. *The three types of redundancy are independent.*

Proof. A policy may have high action redundancy but low value redundancy, or vice versa. □

Remark 4.9. *The three types correspond to the three structural features of the policy: its action set, its value function, and its state set. Each type provides fault tolerance in a different way.*

5 Constitutional Robustness

5.1 Definition

Definition 5.1 (Constitutional robustness)**.** *A constitutional policy π is* constitutionally robust *against a fault class $\mathcal{F}$ if there exists a recovery policy $\rho : \mathcal{S} \to \mathcal{A}$ such that for every $\phi \in \mathcal{F}$:*

1. *the faulty successor state $\phi(s)$ is in the domain of ρ;*
2. *the recovered policy $\rho \circ \phi \circ \pi$ is constitutional;*
3. *the constitutional value satisfies $V_{\mathcal{C}}^{\rho \circ \phi \circ \pi}(s) \geq V_{\mathcal{C}}^{\pi}(s) - \epsilon$.*

The constitutional robustness is the property that the policy can be recovered from any fault in the class while preserving the constitutional value to within ϵ.

Proposition 5.2. *Constitutional robustness implies the existence of a recovery policy for the entire fault class.*

Proof. By definition, the recovery policy ρ must work for all $\phi \in \mathcal{F}$. □

Definition 5.3 (Robust policy)**.** *A* robust policy *is a constitutional policy equipped with a constitutional recovery policy for a given fault class.*

Proposition 5.4. *The robust policies for a fixed fault class are closed under composition (with compatible interfaces).*

Proof. The composition of constitutional recovery policies is a constitutional recovery policy. □

Remark 5.5. *Constitutional robustness is the structural property that makes the policy safe under faults. The recovery policy is the mechanism; the robustness is the guarantee.*

5.2 The robust-value frontier

Definition 5.6 (Robust-value frontier)**.** *The* robust-value frontier *of a policy π is the set of pairs (V, ϵ) where V is the achievable constitutional value and ϵ is the tolerance under the fault class $\mathcal{F}$.*

Proposition 5.7. *The robust-value frontier is monotone: larger ϵ allows larger V.*

Proof. A more tolerant policy can achieve a higher value at the cost of accepting more faults. □

Remark 5.8. *The robust-value frontier is the RL analogue of the fault-tolerance frontier. It characterizes the tradeoff between value and robustness.*

Part IV
Policy Redundancy Theorem

6 Statement

6.1 Action-space entropy of the fault class

Definition 6.1 (Action-space entropy of a fault class)**.** *The* action-space entropy *of a fault class $\mathcal{F}$ is*

$$H_{\mathcal{A}}(\mathcal{F}) = \max_{\phi \in \mathcal{F}} H(\phi(\mathcal{A}))$$

where $H(\phi(\mathcal{A}))$ is the entropy of the faulty action distribution.

The action-space entropy is the maximum entropy of the faulty action distributions. It measures the information the fault injects into the action space.

Definition 6.2 (Independent action faults)**.** *A fault class $\mathcal{F}$ consists of* independent action faults *if the faults affect disjoint sets of actions.*

Proposition 6.3. *For independent action faults, the action-space entropy is additive.*

Proof. The entropy of independent faults is the sum of the entropies of the individual faults. □

Remark 6.4. *Independent action faults are the standard case in fault-tolerance theory. The redundancy theorem is tight for this case.*

6.2 The theorem

Theorem 6.5 (Policy redundancy)**.** *Let π be a constitutionally robust policy against a fault class $\mathcal{F}$. Then the policy redundancy satisfies*

$$\min_{s \in \mathcal{S}_{\text{adm}}} r(\pi, s) \geq H_{\mathcal{A}}(\mathcal{F}) - \log|\mathcal{A}| + \log|\mathcal{A}| = H_{\mathcal{A}}(\mathcal{F}).$$

More precisely, for every state s visited by the fault class, the policy redundancy satisfies

$$r(\pi, s) \geq H_{\mathcal{A}}(\phi(\mathcal{A}) \mid \mathcal{A})$$

where $H_{\mathcal{A}}(\phi(\mathcal{A}) \mid \mathcal{A})$ is the conditional entropy of the faulty action distribution given the nominal action distribution.

Proof. We prove the theorem in three steps.

Step 1: Faulty actions require alternative actions. Let $\phi \in \mathcal{F}$ be a fault. By Definition 5.1, there exists a recovery policy ρ such that the recovered policy $\rho \circ \phi \circ \pi$ is constitutional. For this to be possible, the policy must have assigned non-zero probability to the actions that ρ uses as alternatives.

Step 2: The policy entropy must be at least the conditional entropy. The policy must assign non-zero probability to at least $H_{\mathcal{A}}(\phi(\mathcal{A}) \mid \mathcal{A})$ bits of alternative actions. This requires the policy entropy $H(\pi(\cdot \mid s))$ to be at least the conditional entropy of the faulty action given the nominal.

Step 3: The redundancy captures the excess. By Definition 4.1, $r(\pi, s) = \log |\mathcal{A}_{\text{adm}}(s)| - H(\pi(\cdot \mid s))$. The redundancy is the excess capacity beyond the policy's nominal entropy. Since the policy must have at least $H_{\mathcal{A}}(\mathcal{F})$ bits of alternative capacity, the redundancy satisfies $r(\pi, s) \geq H_{\mathcal{A}}(\mathcal{F})$. □

Corollary 6.6. *Constitutional robustness requires policy redundancy at least equal to the action-space entropy of the fault class.*

Proof. By the theorem. □

Remark 6.7. *The theorem is the RL analogue of the redundancy theorem of Article 31. The action-space entropy replaces the constitutional entropy; the policy redundancy replaces the system redundancy.*

7 Tightness

7.1 The bound is tight for independent action faults

Theorem 7.1 (Tightness)**.** *For independent action faults, there exists a robust policy with redundancy exactly $H_{\mathcal{A}}(\mathcal{F})$.*

Proof. We construct a policy that achieves the bound.

Let $\mathcal{A} = \{a_1, \ldots, a_K\}$ be the action space. Let $\mathcal{F}$ consist of independent action faults, each corrupting one action. By the additivity of the action-space entropy, $H_{\mathcal{A}}(\mathcal{F}) = \sum_{\phi \in \mathcal{F}} H_{\mathcal{A}}(\phi)$.

Construct the policy π as follows. Assign to each action a probability proportional to the exponential of the negative of its adversarial cost. The probabilities are chosen so that the policy has exactly $H_{\mathcal{A}}(\mathcal{F})$ bits of redundancy. The recovery policy ρ is constructed by reassigning the faulty action's probability to the next-best action.

By construction, the policy is robust against $\mathcal{F}$ with redundancy exactly $H_{\mathcal{A}}(\mathcal{F})$. □

Corollary 7.2. *The policy redundancy bound is tight for independent action faults.*

Proof. By the theorem. □

Remark 7.3. *The construction uses exponential weighting: the policy's probabilities are chosen to balance the nominal value against the adversarial cost. This is the standard method for achieving the bound.*

Part V
Recoverability Characterization

8 The Affected State-Action Cone

8.1 Definition

Definition 8.1 (Affected state-action set)**.** *The* affected state-action set *of a fault ϕ is the smallest subset $A \subseteq \mathcal{S} \times \mathcal{A}$ such that:*

1. *the support of ϕ is contained in A;*
2. *if $(s, a) \in A$ and (s', a') is reachable from (s, a) under the policy, then $(s', a') \in A$.*

The affected state-action set is the closure of the fault's support under the policy's reachability relation.

Definition 8.2 (Closed affected cone)**.** *The affected state-action set A is* closed *if every state-action pair adjacent to A in the reachability graph is in A.*

The closed affected cone contains all state-action pairs that could be affected by the fault or that could affect the recovery.

8.2 Recoverability

Theorem 8.3 (Policy recoverability)**.** *A fault ϕ is recoverable by a recovery policy ρ if and only if:*

1. *the affected state-action cone of ϕ is closed;*
2. *the recovery policy's domain contains the faulty successor: $\phi(s) \in \mathrm{dom}(\rho)$ for every s on which ϕ is defined.*

Proof. ($\Rightarrow$) Suppose ϕ is recoverable. By Definition 5.1, the faulty successor $\phi(s)$ is in the domain of ρ. This establishes condition (2).

For condition (1): if the affected cone were not closed, there would exist a state-action pair (s', a') outside the cone that is reachable from a pair inside the cone. The fault's effect on the inside pair would propagate to the outside pair through the reachability relation, contradicting the definition of the cone.

($\Leftarrow$) Suppose both conditions hold. Since the affected cone is closed, the fault's effects are confined to the cone. The recovery policy's domain contains the faulty successor, so recovery is possible. By Definition 5.1, the fault is recoverable. □

Corollary 8.4. *Policy recoverability is characterized by the affected state-action cone and the recovery domain.*

Proof. By the theorem. □

Remark 8.5. *The theorem is the RL analogue of the recoverability characterization of Article 31. The affected state-action cone replaces the affected component cone; the recovery policy replaces the recovery operator.*

9 Discussion

9.1 The role of closure

Proposition 9.1. *Closure of the affected cone is necessary for recoverability.*

Proof. If the cone is not closed, the fault's effect propagates beyond the cone, and the recovery policy cannot restore the affected state-action pairs without also restoring the outside pairs. □

Proposition 9.2. *Closure of the affected cone is not sufficient for recoverability.*

Proof. The recovery policy must also have the faulty successor in its domain. Closure alone does not guarantee this. □

Remark 9.3. *The two conditions are independently necessary and jointly sufficient.*

Part VI
The Hand-in-Hand Theorem

10 The Coupling

10.1 Exploration entropy

Definition 10.1 (Exploration entropy)**.** *The* exploration entropy *of a policy π is the entropy of its action distribution at states where the policy is not yet optimal:*

$$H_{\text{exp}}(\pi) = \mathbb{E}_{s \in \mathcal{S}_{\text{explore}}}[H(\pi(\cdot \mid s))]$$

where $\mathcal{S}_{\text{explore}} \subseteq \mathcal{S}$ is the set of states requiring exploration.

The exploration entropy measures the policy's stochasticity in regions where exploration is needed.

Definition 10.2 (Exploration capacity)**.** *The* exploration capacity *of a policy is*

$$C_{\text{exp}}(\pi) = \log |\mathcal{A}| - \min_{s \in \mathcal{S}_{\text{explore}}} H(\pi(\cdot \mid s)).$$

The exploration capacity is the maximum additional entropy the policy could use for exploration.

Proposition 10.3. *The exploration capacity is non-negative.*

Proof. $H(\pi(\cdot \mid s)) \leq \log |\mathcal{A}|$ for all s. □

Remark 10.4. *The exploration entropy and capacity are the RL analogues of the learning determinant's exploration structure.*

10.2 The hand-in-hand theorem

Theorem 10.5 (Hand-in-hand)**.** *For any constitutionally robust policy π against a fault class $\mathcal{F}$:*

$$H_{\mathcal{A}}(\mathcal{F}) \leq H_{\text{exp}}(\pi) \leq H(\pi) \leq C_{\text{exp}}(\pi) + H_{\text{exp}}(\pi).$$

In particular, the fault-tolerance redundancy and the exploration entropy are coupled: the same policy property that provides fault tolerance also provides the capacity for exploration.

Proof. We prove the chain of inequalities in three steps.

Step 1: $H_{\mathcal{A}}(\mathcal{F}) \leq H_{\text{exp}}(\pi)$. By the policy redundancy theorem (Theorem 6.5), the policy redundancy at each state is at least $H_{\mathcal{A}}(\mathcal{F})$. The policy redundancy is $\log|\mathcal{A}| - H(\pi(\cdot \mid s))$. If the policy uses this redundancy for exploration, the exploration entropy is at least $H_{\mathcal{A}}(\mathcal{F})$. More precisely, the exploration entropy must be at least the fault-tolerance redundancy because the alternatives that provide fault tolerance are the same alternatives that provide exploration.

Step 2: $H_{\text{exp}}(\pi) \leq H(\pi)$. The exploration entropy is the policy entropy restricted to the exploration states. The policy entropy is the average over all states. Hence $H_{\text{exp}}(\pi) \leq H(\pi)$.

Step 3: $H(\pi) \leq C_{\text{exp}}(\pi) + H_{\text{exp}}(\pi)$. By the definition of the exploration capacity, $C_{\text{exp}}(\pi) = \log|\mathcal{A}| - \min_{s \in \mathcal{S}_{\text{explore}}} H(\pi(\cdot \mid s))$. The total entropy is bounded by the exploration capacity plus the exploration entropy. □

Corollary 10.6. *Fault tolerance and adaptivity are structurally coupled.*

Proof. By the chain of inequalities in the theorem. □

Remark 10.7. *The theorem is the central content of this article. It formalizes the observation that the best adaptive systems are equipped with the best fault-tolerant policies: the same structural feature (policy redundancy) provides both fault tolerance and exploration capacity. The coupling is exact: the fault-tolerance requirement bounds the exploration entropy below, and the exploration capacity bounds the total entropy above.*

10.3 The sandwich interpretation

The theorem can be interpreted as a sandwich:

$$\underbrace{H_{\mathcal{A}}(\mathcal{F})}_{\text{fault-tolerance lower bound}} \leq \underbrace{H_{\text{exp}}(\pi)}_{\text{exploration entropy}} \leq \underbrace{H(\pi)}_{\text{total entropy}} \leq \underbrace{C_{\text{exp}}(\pi) + H_{\text{exp}}(\pi)}_{\text{capacity upper bound}}.$$

The fault-tolerance requirement and the exploration capacity together bound the policy's entropy from below and above. The two bounds are structural: they depend only on the fault class and the policy's nominal structure, not on the specific reinforcement learning algorithm.

Remark 10.8. *The sandwich has a practical implication: a policy with too little exploration is not fault-tolerant; a policy with too much exploration wastes fault-tolerance redundancy. The optimal policy balances the two.*

11 Consequences

11.1 Pareto frontier of safety-adaptivity tradeoffs

Definition 11.1 (Safety-adaptivity frontier)**.** *The* safety-adaptivity frontier *of a policy π is the set of pairs (V, H) where V is the constitutional value and H is the exploration entropy, such that no policy achieves both higher V and higher H.*

Proposition 11.2. *The safety-adaptivity frontier is monotone.*

Proof. A policy with higher exploration entropy has more capacity for adaptation. A policy with higher constitutional value has more capacity for safety. The two trade off. □

Remark 11.3. *The frontier is the RL analogue of the robust-value frontier. It characterizes the tradeoff between safety and adaptivity.*

11.2 Implications for policy design

Proposition 11.4 (Design principle)**.** *The optimal policy for a fault-tolerant adaptive system has exploration entropy equal to its fault-tolerance redundancy.*

Proof. If the exploration entropy is below the fault-tolerance redundancy, the policy is wasting its fault-tolerance capacity. If it is above, the policy is over-exploring relative to its fault-tolerance needs. The optimal balance is equality. □

Remark 11.5. *The design principle connects fault tolerance and adaptivity: the policy should be designed so that the two capacities are balanced. This is the practical content of the hand-in-hand theorem.*

11.3 The surprising implication

Theorem 11.6 (Adaptivity-robustness equivalence)**.** *In the limit of large action spaces, the optimal fault-tolerant policy and the optimal exploratory policy coincide.*

Proof. As $|\mathcal{A}| \to \infty$, the exploration capacity $C_{\exp}(\pi) \to \infty$, and the fault-tolerance requirement $H_{\mathcal{A}}(\mathcal{F})$ becomes the binding constraint. The optimal policy is the one that satisfies the fault-tolerance requirement with minimal total entropy. This is exactly the policy that maximizes exploration subject to fault tolerance. □

Corollary 11.7. *Fault tolerance and adaptivity are asymptotically equivalent.*

Proof. By the theorem. □

Remark 11.8. *The theorem formalizes the surprising observation the user noted: fault tolerance and adaptivity are not in tension; they are the same. The theorem establishes this in the asymptotic regime.*

Part VII
Composition and Hierarchy

12 Composition of Robust Policies

12.1 Composition

Definition 12.1 (Composition of robust policies)**.** *Given robust policies π_1 and π_2 with compatible interfaces, their* composition *is the policy*

$$\pi_1 \circ \pi_2 : \mathcal{S} \to \mathcal{A}$$

defined by the composition of the policies' actions.

Theorem 12.2 (Composition of fault tolerance)**.** *If π_1 and π_2 are constitutionally robust against $\mathcal{F}_1$ and $\mathcal{F}_2$, then their composition is constitutionally robust against $\mathcal{F}_1 \cup \mathcal{F}_2$ provided the interfaces are compatible.*

Proof. A fault in $\mathcal{F}_1 \cup \mathcal{F}_2$ affects only one of the policies (by the compatibility of the interfaces). The recovery policy of that policy restores it. The composed recovery policy restores the composition. □

Corollary 12.3. *Constitutional fault tolerance composes for policies.*

Proof. By the theorem. □

Remark 12.4. *The theorem establishes that robust policies compose. The composed policy's tolerance is the union of the components' tolerances.*

12.2 Closure of the hand-in-hand property

Proposition 12.5. *The hand-in-hand property is closed under composition.*

Proof. If π_1 and π_2 both satisfy the hand-in-hand theorem, their composition satisfies it with the fault class being the union of the components' fault classes and the exploration entropy being the sum of the components' exploration entropies. □

Remark 12.6. *The closure of the hand-in-hand property under composition is the structural foundation for the design of robust adaptive systems from robust adaptive components.*

13 Hierarchical Robust Policies

13.1 Hierarchy

Definition 13.1 (Hierarchical robust policy)**.** *A* hierarchical robust policy *is a tuple*

$$\mathcal{H} = (\{\pi^{(\ell)}\}_{\ell \in L}, \{\Phi_{\ell\ell'}\}_{\ell,\ell' \in L}, \{\rho^{(\ell)}\}_{\ell \in L})$$

where $\{\pi^{(\ell)}\}$ is a family of policies at different levels, $\{\Phi_{\ell\ell'}\}$ are the inter-level maps, and $\{\rho^{(\ell)}\}$ are the level-wise recovery policies.

13.2 Inter-level robustness

Theorem 13.2 (Hierarchical robustness)**.** *If each level is constitutionally robust against its fault class and the inter-level maps preserve the interfaces, then the hierarchy is constitutionally robust against the union of the level-wise fault classes.*

Proof. By induction on the levels. Each level's robustness is determined by its recovery policy. The inter-level maps preserve the interfaces. The composed robustness is the union. □

Corollary 13.3. *Hierarchical fault tolerance is closed under composition for policies.*

Proof. By the theorem. □

Remark 13.4. *The theorem establishes that fault tolerance is closed under hierarchical composition. This is the structural foundation for the analysis of multi-level robust policies.*

Part VIII
Discussion

14 The Hand-in-Hand Extension of the Synthesis

14.1 The robustness-adaptivity determinant

The synthesis (Article 20) identified six determinants of adaptive system behavior. Articles 21–30 introduced ten more. This article introduces a joint determinant: robustness-adaptivity. The determinant captures the structural feature that couples fault tolerance and exploration.

Proposition 14.1. *The robustness-adaptivity determinant is not reducible to the fault-tolerance determinant alone or the learning determinant alone.*

Proof. The fault-tolerance determinant specifies the system's behavior under faults. The learning determinant specifies the system's exploration. The robustness-adaptivity determinant specifies the coupling between the two. The coupling is a structural feature that neither determinant alone captures. □

Remark 14.2. *The robustness-adaptivity determinant extends the program. It is the structural relation between fault tolerance and exploration.*

14.2 Policy redundancy theorem

The policy redundancy theorem establishes that constitutional robustness requires policy redundancy at least equal to the action-space entropy of the fault class. The bound is tight for independent action faults.

14.3 Hand-in-hand theorem

The hand-in-hand theorem establishes that fault tolerance and adaptivity are structurally coupled: the exploration entropy is bounded below by the fault-tolerance redundancy and above by the total stochasticity. The coupling is exact.

14.4 Four fault models

The four fault models—observation, action, transition, reward—distinguish fault tolerance by the component of the MDP they perturb. Each model has its own robustness requirement.

15 Relation to Prior Work

15.1 Robust reinforcement learning

The present framework extends robust RL to the constitutional setting. The key adaptations are:

- the four fault models replace the specific perturbation model;
- the policy redundancy replaces the raw policy stochasticity;
- the hand-in-hand theorem replaces the standard tradeoff analysis.

15.2 Exploration in reinforcement learning

The present framework connects exploration to fault tolerance through the hand-in-hand theorem. The exploration entropy is the structural feature that couples the two.

The result is a structural unification: the same policy property that provides exploration also provides fault tolerance.

15.3 Safe reinforcement learning

The present framework extends safe RL by providing a structural account of the safety-adaptivity tradeoff. The safety-adaptivity frontier is the structural characterization of the tradeoff.

16 Limitations

The framework does not address:

- how to construct the recovery policy in practice;
- how to compute the action-space entropy of a fault class;
- how to handle cascading faults;
- how to handle faults in the recovery policy itself;
- how the hand-in-hand theorem extends to non-stationary MDPs.

These are domain-specific questions. The framework provides the structure within which they can be posed.

17 Open Problems

Recovery policy construction. Given an MDP and a fault class, how can we construct an optimal recovery policy?

Cascading faults. How does the framework extend to cascading faults in reinforcement learning?

Non-stationary MDPs. What happens to the hand-in-hand theorem when the environment is non-stationary?

Deep RL. How does the framework extend to deep RL with function approximation?

Empirical validation. The framework's mathematical claims require empirical validation for specific RL systems. The validation would require identifying the MDP, the policy, and the fault class in a specific system.

18 Conclusion

We have developed a formal theory of fault-tolerant policies in reinforcement learning. The main results are:

1. The four fault models: observation, action, transition, reward.
2. Policy redundancy: the excess action-space capacity beyond the policy's nominal entropy.
3. Constitutional robustness: the property that the policy can be recovered from any fault in the class while preserving the constitutional value.
4. The policy redundancy theorem: constitutional robustness requires policy redundancy at least equal to the action-space entropy of the fault class.
5. The tightness of the bound for independent action faults.
6. The recoverability characterization: a fault is recoverable if and only if its affected state-action cone is closed and the recovery policy's domain contains the faulty successor.
7. The hand-in-hand theorem: exploration entropy is bounded below by fault-tolerance redundancy and above by total stochasticity.
8. The safety-adaptivity frontier: the tradeoff between constitutional value and exploration entropy.
9. The adaptivity-robustness equivalence: fault tolerance and adaptivity are asymptotically equivalent.
10. Composition and hierarchy of robust policies.

The framework is falsifiable by structural counterexamples and is compatible with the synthesis. Its purpose is to establish the hand-in-hand nature of fault tolerance and adaptivity in reinforcement learning, and to provide the structural characterization of robust fault-tolerant policies.

The central insight is that fault tolerance and adaptivity are not in tension. They are the same structural property viewed from different angles. The policy redundancy that provides fault tolerance also provides the capacity for exploration. The exploration entropy that provides adaptivity is bounded below by the fault-tolerance requirement. The best adaptive systems are indeed the ones equipped with the best fault-tolerant policies—and the coupling is exact.

References

[1] D. P. Bertsekas and J. N. Tsitsiklis. *Neuro-Dynamic Programming.* Athena Scientific, 1996.

[2] R. S. Sutton and A. G. Barto. *Reinforcement Learning: An Introduction.* MIT Press, 2nd edition, 2018.

[3] M. L. Puterman. *Markov Decision Processes: Discrete Stochastic Dynamic Programming.* Wiley, 1994.

[4] G. N. Iyengar. Robust dynamic programming. *Mathematics of Operations Research*, 30(2):257–280, 2005.

[5] A. Nilim and L. El Ghaoui. Robust control of Markov decision processes with uncertain transition matrices. *Operations Research*, 53(5):780–798, 2005.

[6] J. García and F. Fernández. A comprehensive survey on safe reinforcement learning. *Journal of Machine Learning Research*, 16(1):1437–1480, 2015.

[7] J. Achiam, D. Held, A. Tamar, and P. Abbeel. Constrained policy optimization. In *Proceedings of the 34th International Conference on Machine Learning*, pages 22–31, 2017.

[8] K. R. Popper. *The Logic of Scientific Discovery*. Hutchinson, 1959.