

The Six-Layer Linguistic Framework

A Standard and Operational Handbook for Prose Elevation

Version 1.0 · June 2026 · issued for routine use

Contents

1	Purpose, Scope, and Operating Contract	2
1.1	Objective	2
1.2	Input and output	2
1.3	Operating modes	2
1.4	Scope and limits	2
2	The Six Layers as Measures of Quality	2
2.1	L1 — Lexical–Syntactic (the sentence)	3
2.2	L2 — Structural–Organizational (the arrangement)	4
2.3	L3 — Semantic–Entity (reference)	4
2.4	L4 — Discourse–Pragmatic (flow)	5
2.5	L5 — Argumentative–Rhetorical (earning conclusions)	6
2.6	L6 — Meta-Textual (the whole, for its purpose)	7
3	Operating Procedure	8
4	Fidelity: The One Boundary	9
5	Scoring, Disposition, and the Holistic Gate	9
6	Reporting: State Each Point Once	10
7	Appendix A — Quick Reference	10
8	Appendix B — Glossary	11

Input: any prose document, in any language. *Output:* the same document elevated to the best state it can reach without betraying its origin, returned as tracked changes, accompanied by a short report. This handbook is written to be read by a human editor and executed directly by a language model.

1 Purpose, Scope, and Operating Contract

1.1 Objective

This framework raises prose to the best version of itself. Its operative question is not “*where does this text break a rule?*” but “*what is the best form of this passage that a skilled editor could produce without changing what it means or the author’s position?*” The six layers of Section 2 supply the definition of *better*; the fidelity boundary of Section 4 supplies the one limit on how far the editor may go. Everything else here serves those two ideas.

The default action is therefore **rewrite, not flag**. Every passage is rewritten to the strongest form it can reach. A passage is left as written only when it is already *exemplary* — when no rewriting could make it materially better within the one boundary of Section 4. The target is the best the text can become, not a tidied copy of what arrived. **Minimal, cosmetic, or near-invisible revision is itself a failure of this framework**: if the elevated text is not visibly and substantially stronger than the input, the framework has not been applied. The instrument is elevation — not a defect checklist, and not a light copyedit.

1.2 Input and output

The framework accepts any document from which running text can be read, at any stage of development. It returns two things and only two things: the **elevated document, carrying every edit as a tracked change with a brief rationale**, and a **lean report** (Section 6). There are no parallel logs.

1.3 Operating modes

Audit. Evaluate, score, and report without rewriting. Appropriate for triage, external review, and acceptance testing.

Elevate (default). Apply the full procedure and return the elevated document and the report. Unless a caller states otherwise, Elevate is assumed throughout.

1.4 Scope and limits

The framework applies to expository, technical, policy, and argumentative prose of any length and any language. It is not built for verse, fiction, or short-form marketing copy, where deliberate violation of these measures is often the technique. The measures are judgments, not arithmetic: a construction that demonstrably serves the work is kept, whatever a generic measure would say. Reports cite the version of this handbook under which they were produced.

2 The Six Layers as Measures of Quality

The framework reads prose through six layers, ordered from the most local linguistic unit to the most global. Each is a **measure of what *better* looks like**, not a checklist of violations: each states an objective, the measures that define excellence at that layer, and the moves that carry prose toward it. The editor applies every layer’s measures; how, and in what order, is governed by Section 3. The measures are **size-independent**: rate properties are stated as ratios or densities, and singular properties as presence or per-instance, so the same measures govern a memo and a monograph alike (Section 5).

2.1 L1 — Lexical–Syntactic (the sentence)

Better means every sentence is correct; each concept keeps one settled designation; agency is carried directly, so the reader sees who does what; sentence length and opening vary; and no word is idle. *Moves*: recast for directness, tighten, de-nominalize, vary cadence, and normalize terminology — while leaving terms of art and the author’s characteristic phrasing intact.

Code	Indicator and default threshold	Default severity
L1.1	Grammatical correctness. Sentences contain no errors of agreement, tense, case, or punctuation that a standard grammar would flag.	Minor; Major where the error obscures meaning.
L1.2	Pronoun reference. Every third-person pronoun and demonstrative resolves to exactly one antecedent within the same or immediately preceding sentence.	Major in claim-bearing sentences; otherwise Minor.
L1.3	Terminology consistency. Each concept carries one approved designation per the term registry; variants are not introduced without justification.	Major for key concepts; Minor for peripheral terms.
L1.4	Sentence variety. Mean sentence length within 15–25 words; standard deviation of at least 6; no run of four or more sentences sharing the same opening structure.	Minor.
L1.5	Voice. Passive constructions account for no more than 25% of finite clauses (genre-adjustable); agency is explicit in claim-bearing sentences.	Major where passive voice conceals a responsible agent in a claim; otherwise Minor.
L1.6	Economy. No non-technical content word recurs more than three times per 100 words; no assertion is repeated verbatim without rhetorical purpose.	Minor; Major where repetition pads a load-bearing section.

2.2 L2 — Structural–Organizational (the arrangement)

Better means every section has one clear job; the reader can predict where information lives; units open by declaring their point; what the opening promises the body delivers, and the close adds nothing new; and length tracks importance. *Moves*: reorder, merge, split, rebalance, and add signposting.

Code	Indicator and default threshold	Default severity
L2.1	Section purpose. Every section maps to exactly one declared function (orientation, development, synthesis, or apparatus).	Major.
L2.2	Topic sentences. At least 80% of body paragraphs open with a topic sentence within their first two sentences.	Minor; Major in sections carrying the central argument.
L2.3	Sequence logic. Section order follows one declared scheme: general-to-specific, problem–solution, chronological, or the genre-fixed order.	Major; Critical where disorder prevents navigation.
L2.4	Alignment. Every commitment made in the introduction is met in the body; the conclusion introduces no new substantive material.	Major; Critical where the central commitment is unmet.
L2.5	Proportion. Body sections sit within a 3:1 length ratio relative to functional importance; paragraphs run 3–8 sentences in the typical case.	Minor; Major where a core section is underdeveloped.

2.3 L3 — Semantic–Entity (reference)

Better means every paragraph is traceable to the document’s scope; entities are introduced in full before reduced or pronominal forms; reference chains stay unambiguous and are re-anchored after a gap; vocabulary within a topic holds to one field and register; and no attribute ascribed to an entity contradicts another. *Moves*: repair introductions, re-anchor chains, normalize vocabulary. A genuine conflict of fact is flagged, not invented away (Section 4).

Code	Indicator and default threshold	Default severity
L3.1	Topic anchoring. Every paragraph is traceable to the document's stated scope.	Major; Critical for sustained drift across a section.
L3.2	Entity introduction. Every entity appears in full designation, with abbreviations expanded, before any reduced or pronominal reference.	Major.
L3.3	Coreference continuity. Reference chains are re-anchored with the full designation after a gap of more than one paragraph and at section boundaries.	Minor; Major where the gap creates ambiguity.
L3.4	Semantic-field coherence. Vocabulary within a topic is drawn from a consistent semantic field and register.	Minor.
L3.5	Topic transition. Adjacent paragraphs share at least one referential or lexical link; topic shifts are announced.	Major.
L3.6	Entity consistency. Attributes, values, and roles ascribed to an entity do not conflict anywhere in the document.	Critical.

2.4 L4 — Discourse–Pragmatic (flow)

Better means each unit ties to its neighbours; every connective names the relation that actually holds; markers are not overused; key terms recur without unmotivated breaks; and no paragraph is orphaned. *Moves*: stitch openings, replace or remove inaccurate markers, reinstate broken lexical chains, integrate or relocate orphans.

Code	Indicator and default threshold	Default severity
L4.1	Connected openings. Each body paragraph opens with a tie to prior text: a referential link, a repeated key term, or a discourse marker.	Minor; Major at section boundaries.

Code	Indicator and default threshold	Default severity
L4.2	Marker accuracy. Every discourse marker matches the logical relation it asserts (contrast, cause, addition, concession).	Major; an inaccurate marker asserts a false relation.
L4.3	Marker economy. Sentence-initial markers appear in no more than one of every two to three sentences.	Minor.
L4.4	Lexical chains. The document's key terms recur across sections without unmotivated breaks.	Minor; Major where a central concept's chain breaks.
L4.5	Connectivity. Every paragraph except the first and last has at least one inbound and one outbound link to its neighbours.	Major; an orphan paragraph is unintegrated information.

2.5 L5 — Argumentative–Rhetorical (earning conclusions)

Better means every claim is explicit, locatable, and scoped; load-bearing claims are supported; the step from evidence to claim is visible where it is not self-evident; the strongest credible counter-position is met for contestable central claims; support contains no circular or forward-dependent links; and claim strength matches the strength of the evidence the author already provides. *Moves*: surface buried claims, make warrants explicit, reorder support, and calibrate modality toward the author's own evidence. The editor never manufactures evidence; an unsupported claim is calibrated or flagged, never propped up with invented support.

Code	Indicator and default threshold	Default severity
L5.1	Claim explicitness. Every claim is stated, locatable, and scoped; the central claim appears in the introduction or its genre-fixed position.	Critical for an implicit central claim; Major for supporting claims.

Code	Indicator and default threshold	Default severity
L5.2	Evidence sufficiency. Every load-bearing claim is supported by at least one evidence unit; the central claim by at least two independent units. No section is claim-dense but evidence-light.	Critical for the central claim; Major otherwise.
L5.3	Warrant presence. Where the step from evidence to claim is not self-evident, the warrant is stated.	Major for contestable links.
L5.4	Counterargument. The strongest credible alternative position is addressed for every contestable central claim.	Major.
L5.5	Progression. No claim depends on a claim established later; support relations contain no cycles.	Critical for circular support; Major for forward dependency.
L5.6	Calibration. Claim strength matches evidence strength: hedging is neither absent where evidence is weak nor excessive where it is strong.	Major for overclaiming; Minor for overhedging.

2.6 L6 — Meta-Textual (the whole, for its purpose)

Better means conceptual and lexical complexity match the declared audience; stance and register hold steady; the obligatory moves of the genre are present; every section earns its place against the purpose; and, on a continuous read, the document achieves what it set out to do. *Moves*: tune complexity to the audience, normalize stance, supply a missing obligatory move, and cut or justify a section that does not serve the purpose.

Code	Indicator and default threshold	Default severity
L6.1	Audience fit. Lexical and conceptual complexity matches the declared audience; assumed knowledge is audited against the profile.	Major; Critical where the audience cannot use the document.

Code	Indicator and default threshold	Default severity
L6.2	Stance consistency. Person, register, and hedging level remain stable document-wide.	Minor; Major where drift undermines authority.
L6.3	Genre conformity. All obligatory moves of the declared genre are present (for example: executive summary, methods, recommendations).	Major per missing obligatory move.
L6.4	Purpose contribution. Every section is justified against the purpose statement.	Major; unjustified sections dilute the document.
L6.5	Holistic effectiveness. On a continuous read-through, the document achieves its declared purpose for its declared audience.	Critical; this indicator gates release.

3 Operating Procedure

The procedure fixes the **measures, not a ritual**. What is mandatory is that every layer’s measures in Section 2 are applied and demonstrably present in the result. What is *not* mandated is how many times the text is read, or in what order the layers are worked.

A competent editor — human or agent — moves through three movements, in as few passes as the document needs:

1. **Orient**. Establish the purpose, audience, and genre, and mark the fidelity invariants in the source: the claims, facts, figures, quotations, citations, and the stance to be preserved (Section 4).
2. **Elevate**. Rewrite each passage to the strongest form it can reach against the six layers’ measures, recording each edit as a tracked change with a one-line rationale. Expect substantial, visible change: reaching the best version is the task, not preserving the shape of the original. Often this is a single diagnose-and-elevate pass, returning to a span only when a later change disturbs an earlier one.
3. **Attest**. Confirm that the fidelity boundary held, that the author’s stance survived, that no edit introduced a new defect, and that the document passes the holistic gate (Section 5).

An optional method. Working the layers *countercurrent* — from L6 down to L1, settling purpose, structure, and argument before local polish — is offered because large moves can otherwise waste sentence-level work already done. It is a recommended habit with a clear rationale, **not a required order**. The one durable rule is practical, not ceremonial: do not polish what you may later cut.

Depth calibration. Elevation works at the level of the passage, not the typo. For the expected *magnitude*: a sentence such as “There is a tendency on the part of users to make use of the system

in those cases where it is required” should return as “Users use the system when they must” — shorter, sharper, identical in meaning. If a paragraph comes back with one comma moved and a synonym swapped, it has not been elevated; it has been dusted.

4 Fidelity: The One Boundary

Elevation is bounded by exactly one thing. The governing principle is **free over form, faithful over essence**.

Free to change — form. Sentence structure and syntax; non-terminological word choice; paragraph boundaries and shape; transitions; the order of material within a unit where logic permits; concision; and rhythm. Over these, the editor has full latitude.

Never change — essence and integrity. The meaning and every claim; the author’s positions and conclusions; facts, numbers, dates, names, quotations, and citations; and the author’s stance and register. Nothing that alters *what is asserted*, *what is meant*, or *the author’s position* may be changed. How it is said — every sentence, every structure, every cadence — is open.

Flag, do not fix. When a passage reads as a factual error, an internal contradiction, or a claim that outruns the author’s own evidence, the editor flags it for the author. It is never silently rewritten, and it is never “corrected” with content the editor supplies. Repairing substance is the author’s prerogative; the editor’s reach ends at form.

Voice protects identity, not wording. The text should still read as this author thinking — same stance, same register — not as a different writer. But voice is never a reason to keep a weaker sentence: phrasing, structure, and rhythm are rewritten freely and fully. Preserving the author’s identity and rewriting the prose to its best are not in tension; conflating them is precisely what produces timid, near-invisible edits.

5 Scoring, Disposition, and the Holistic Gate

Each layer is scored on a 1–5 anchored scale, recorded **before and after** so the gain is visible.

	Score	Anchor
5 — Exemplary		The layer actively strengthens the document; nothing material is left to improve.
4 — Proficient		Strong; only minor, isolated room remains.
3 — Adequate		Sound but clearly improvable.
2 — Deficient		Material weaknesses concentrated or repeated.
1 — Failing		Weakness pervades the layer and undermines the document.

Size independence. Scoring rests on the *severity and spread* of weaknesses, never on their raw count, so it does not move with document length. A rate weakness — say, pervasive passive voice — is judged by its proportion; a singular weakness — an unsupported central claim, a contradiction — by its presence and reach. The one qualitative exception is a very short text, where a rate computed over few sentences is statistically noisy and is therefore read by judgment rather than by number.

Disposition. The release bar is **Exemplary**. *Release* only when every layer reaches 5 — actively strengthening the document, with nothing material left to improve — the holistic gate is met, and no flagged item blocks use. A layer at 4 (Proficient) is not finished; it returns to elevation. *Conditional release* when the text is exemplary but flagged substantive items await the author’s decision; these travel with the document. *Recycle* only when a layer cannot reach 5 without breaching the Section 4 boundary — the limit, and the reason, are recorded.

The holistic gate. Regardless of layer scores, the document is read once, whole, against its purpose and audience. If it does not achieve its purpose on that read, it is not released.

6 Reporting: State Each Point Once

The record of the work is the **tracked-changes document itself**. Edits are not retyped into the report. The report repeats nothing that the tracked changes already show, and it contains exactly three parts:

1. **Scorecard.** One line per layer: initial score → final score. No prose.
2. **Summary of work.** A brief account of the *kinds* of elevation performed, with a few representative before/after examples — not an exhaustive list. The full diff is in the document.
3. **For the author’s decision.** A short list of the only items that need human judgment: the flagged factual issues and contradictions from Section 4, and any genuine editorial judgment calls.

Parallel logs, cross-referenced finding-and-change registers, and re-enumeration of the same edit in several places are prohibited. Each fact appears once.

7 Appendix A — Quick Reference

Layer	The measure of <i>better</i>
L1 Lexical–Syntactic	Correct, direct, varied, economical sentences.
L2 Structural	One job per section; predictable, proportioned arrangement.
L3 Semantic–Entity	Stable, unambiguous reference; consistent field; no contradictions.
L4 Discourse	Connected units; accurate, unpadding markers; unbroken key-term chains.
L5 Argumentative	Explicit claims, visible warrants, met counter-positions, calibrated strength.
L6 Meta-Textual	Right for its audience, genre, and purpose on a whole read.

Contract: Rewrite every passage to *exemplary* — minimal or invisible change is a failure; release only when every layer reaches 5. **Boundary:** free over form, faithful over essence; flag, don’t fix. **Procedure:** apply every measure, in as few passes as needed; countercurrent optional. **Report:** tracked changes are the record; scorecard + summary + author’s-decision list; each point once.

8 Appendix B — Glossary

Elevation. Rewriting a passage to the strongest form it can reach under the fidelity boundary — to exemplary, not merely improved; the default action of this framework.

Fidelity. The single limit on elevation: faithfulness to the original’s essence and integrity — meaning, claims, facts, positions, and stance.

Essence and integrity. What may not change: what the text asserts, what it means, the author’s position, and stance.

Form. What may change freely: syntax, non-terminological wording, paragraph shape, transitions, internal order, concision, and rhythm.

Measure. A statement of what excellence looks like at a layer, used to guide rewriting — as opposed to a threshold used to detect a violation.

Flag, do not fix. The rule that suspected errors of substance are surfaced for the author rather than silently rewritten.

Voice. The author’s stance and register, preserved as identity — never invoked as a reason to keep weaker wording.